



Current experiment architecture

Both datasets share the same search and validation procedure. Frozen inputs are read; experiment results are written to a separate run directory.



The same prediction + evaluation engine runs at every stage below; only partition IDs and candidate settings change.



Frozen grouped dataset: base.json
720 questions; 432 / 144 / 144. Same non-empty gold set stays in one split.

Reviewed dataset: manual_review_v1.json
710 questions; 426 / 142 / 142. Question-level random split, seed 42.